

Developing a Large Language Model-Based Playlist Continuation Recommender System

Enzo CHAROLOIS–PASQUA, Eléa VELLARD

Under the supervision of
Youssra REBBOUD, Pasquale LISENA and Raphael TRONCY

Semester project, January 2025



Table of Contents

- ① Introduction
- ② Related Work
- ③ Method
- ④ Fine-tuning the model
- ⑤ Playlist Generation
- ⑥ Results
- ⑦ Conclusion and Future Work

Table of Contents

- ① Introduction
- ② Related Work
- ③ Method
- ④ Fine-tuning the model
- ⑤ Playlist Generation
- ⑥ Results
- ⑦ Conclusion and Future Work

Introduction

Context:

- Music streaming platforms (e.g., Spotify) play a key role.
- Personalized recommendation is essential: automatically generating tailored playlists.
- This project builds upon the RecSys Challenge 2018 and the Million Playlist Dataset (MPD).

Overall Goal:

- Explore an up-to-date playlist generation approach using LLMs.
- Enhance playlist creation with semantic clustering, embeddings, and fine-tuning of a sentence-transformer model.

Table of Contents

① Introduction

② Related Work

③ Method

④ Fine-tuning the model

⑤ Playlist Generation

⑥ Results

⑦ Conclusion and Future Work

Related Work

Recommender systems and music recommendation:

- Robin Burke, Alexander Felfernig, and Mehmet H Göker (2011). “Recommender systems: An overview”. In: *Ai Magazine* 32.3, pp. 13–18
- Ching-Wei Chen et al. (2018). “Recsys challenge 2018: Automatic music playlist continuation”. In: *Proceedings of the 12th ACM Conference on Recommender Systems*, pp. 527–528

Previous approaches:

- Diego Monti et al. (2018). “An ensemble approach of recurrent neural networks using pre-trained embeddings for playlist completion”. In: *Proceedings of the ACM Recommender Systems Challenge 2018*, pp. 1–6

Technical assets:

- Nils Reimers and Iryna Gurevych (Nov. 2019). “Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks”. In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing*. Association for Computational Linguistics, pp. 3982–3992
- Guy Shani and Asela Gunawardana (2011). “Evaluating Recommendation Systems”. In: *Recommender Systems Handbook*. Ed. by Francesco Ricci et al. Boston, MA: Springer US, pp. 257–297

Table of Contents

① Introduction

② Related Work

③ Method

④ Fine-tuning the model

⑤ Playlist Generation

⑥ Results

⑦ Conclusion and Future Work

Dataset

Spotify's Million Playlist Dataset

- 1,000,000 user-generated playlists.
- January 2010 - October 2017.
- 1,000 JSON slices.

Method Overview

Main Steps:

- ① **MPD Preprocessing:** convert JSON to CSV, separating playlists/items/tracks.

Method Overview

Main Steps:

- ① **MPD Preprocessing:** convert JSON to CSV, separating playlists/items/tracks.
- ② **Semantic clustering:** embeddings and K-Means algorithm.

Method Overview

Main Steps:

- ① **MPD Preprocessing:** convert JSON to CSV, separating playlists/items/tracks.
- ② **Semantic clustering:** embeddings and K-Means algorithm.
- ③ **Fine-tuning a lightweight transformer:** classify playlists into thematic clusters.

Method Overview

Main Steps:

- ① **MPD Preprocessing:** convert JSON to CSV, separating playlists/items/tracks.
- ② **Semantic clustering:** embeddings and K-Means algorithm.
- ③ **Fine-tuning a lightweight transformer:** classify playlists into thematic clusters.
- ④ **Playlist Generation:** voting system.

Method Overview

Main Steps:

- ① **MPD Preprocessing:** convert JSON to CSV, separating playlists/items/tracks.
- ② **Semantic clustering:** embeddings and K-Means algorithm.
- ③ **Fine-tuning a lightweight transformer:** classify playlists into thematic clusters.
- ④ **Playlist Generation:** voting system.
- ⑤ **Evaluation:** Quantitative and qualitative measures.

Clustering: Embedding Model (1/3)

Idea: Identify semantic themes among playlists¹.

- FastText: lightweight but did not capture enough the contextual meaning of playlist titles.

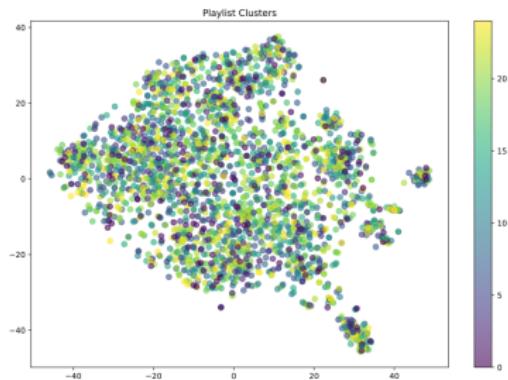


Figure 1: T-SNE plot of training data usign FastText

¹Bohan Li et al. (2020). "On the sentence embeddings from pre-trained language models". In: *arXiv preprint arXiv:2011.05864*.

Clustering: Embedding Model (2/3)

- **BERT (Bidirectional Encoder Representations from Transformers):**
 - Uses self-attention mechanisms to learn contextual representations of tokens.
 - Not optimized to produce one fixed embedding for an entire sentence.
- **Sentence-BERT:**
 - Extends BERT with a *Siamese* architecture: two BERT modules process sentence pairs in parallel.²
 - Adds a pooling layer on top to produce fixed-size sentence embeddings.
 - Specially designed for sentence-level tasks.

Why Sentence-BERT for our project?

- More effective at capturing the semantic meaning of playlist titles.
- Efficient similarity comparison among a large set of sentences or playlists.

²Nils Reimers and Iryna Gurevych (Nov. 2019). "Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks". In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing*. Association for Computational Linguistics, pp. 3982–3992.

Clustering: Embedding Model (3/3)

① **Model selection:** Sentence-BERT

② **Content-based approach:** The playlist embedding is computed as the average of its track embeddings:

$$\mathbf{e}_{\text{playlist}} = \frac{1}{n} \sum_{i=1}^n \mathbf{e}_{\text{track}_i}$$

③ **Resulting Representation:** The resulting vector represents the overall semantic theme of the playlist.

Why Averaging?

- Averaging helps to smooth out noise and variability among individual tracks.
- It provides a stable, fixed-size representation for each playlist.

Clustering: Algorithm (1/2)

Why K-Means?

- A simple, well-known clustering algorithm.
- It partitions data into a predefined number of clusters by minimizing intra-cluster variance.

Why not DBScan:

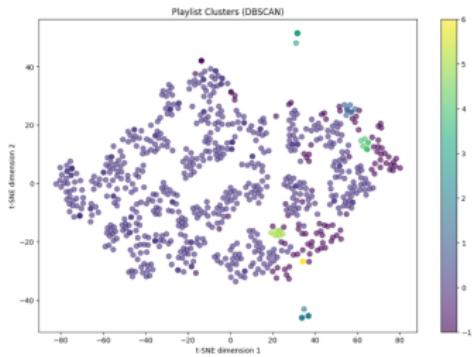


Figure 2: T-SNE plot of training data using DBScan

Clustering: Algorithm (2/2)

Choosing k :

- We empirically set $k = 200$ to maximize intra-cluster coherence.
- Other techniques (like the elbow method) were used to guide the choice, but our focus was on cluster coherence.

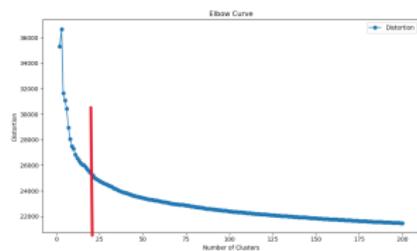


Figure 3: Elbow curve method. Suggesting an optimal $k \approx 22$.

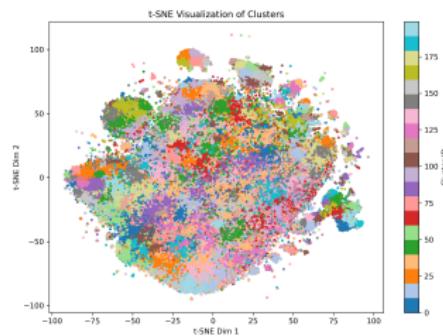


Figure 4: T-SNE plot for the whole dataset, using Sentence-BERT and K-Means algorithm.

Clustering: cleaning

Miscellaneous Clusters:

- Some clusters tend to be very large and less coherent.

Cleaning Strategy:

- **Post-Clustering Analysis:** Evaluate clusters for internal coherence.
Simulate a manual check using the percentage of exact title matches.
- **Filtering:** Cleaning threshold $t = 2$.

Clustering: results

- **Cleaning results:** removed around 40% of playlists, of 21% clusters.
- **Final k:** $k = 158$
- **Split:** (80/10/10) split, ensuring a representation of each cluster in both sets.

Representation of each cluster in split

- So the model **learns** and is **evaluated** on the full diversity of playlists.
- To avoid **bias** towards clusters that might be overrepresented in one split.

Table of Contents

- ① Introduction
- ② Related Work
- ③ Method
- ④ Fine-tuning the model
- ⑤ Playlist Generation
- ⑥ Results
- ⑦ Conclusion and Future Work

Fine-tuning the model: Goal

Idea: Adapt a pre-trained sentence-transformer model to better capture the semantic nuances of playlist titles,

- **Data:** Playlists titles
- **Labels:** Cluster IDs
- Using the **training set** to update its weights and adjust model parameters for improved thematic classification.
- Using the **validation set** to monitor performance and prevent overfitting.

Fine-tuning the model: Architecture (1/2)

Base Model: *all-MiniLM-L6-v2*

- A compact sentence Transformers, providing 384-dimensional embeddings³⁴⁵.
- Adapted for classification.

Classification head:

- A dropout layer to prevent overfitting.
- A fully connected layer mapping the 384-dimension to the number of clusters.
- A mapping from original cluster IDs to sequential labels

³Omar Espejel (2022). "Train and Fine-Tune Sentence Transformers Models". In: *Hugging Face*.

⁴Yichu Zhou and Vivek Srikumar (2021). "A Closer Look at How Fine-tuning Changes BERT". In: *CoRR* abs/2106.14282. arXiv: 2106.14282. URL: <https://arxiv.org/abs/2106.14282>.

⁵Chi Sun et al. (2019). "How to fine-tune bert for text classification?" In: *Chinese computational linguistics: 18th China national conference, CCL 2019, Kunming, China, October 18–20, 2019, proceedings 18*. Springer, pp. 194–206.

Fine-tuning the model: Architecture (2/2)

Training Setup:

- **Loss Function:** Cross-entropy loss, standard for classification tasks.
- **Batch Size:** 8 (computation constraints).
- **Epochs:** 5 (computation constraints).
- **Learning Rate:** 2×10^{-5} to prevent overfitting (higher generalization accuracy).

Evaluation Strategy:

- Evaluation after each epoch.
- Monitoring both training and validation loss.
- Saving the model with the best validation performance.

Fine-tuning the model: Results

Observations:

- A clear decrease in loss.
- No signs of overfitting.

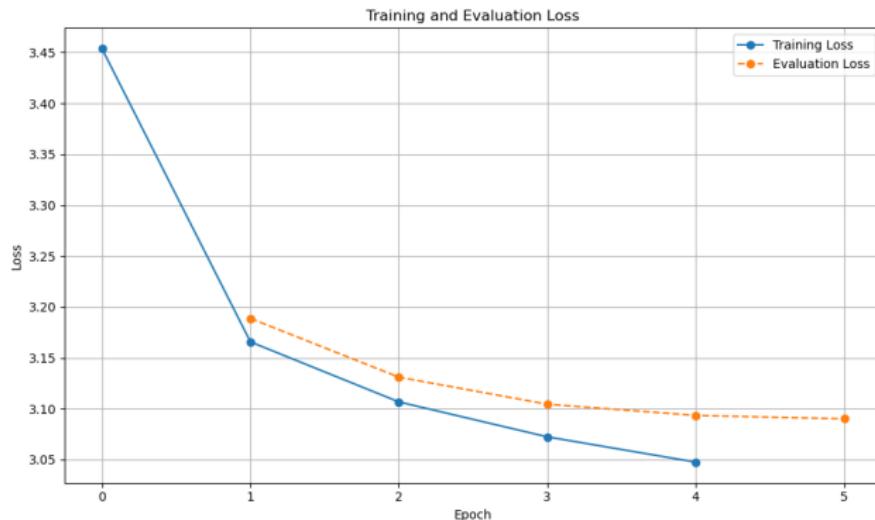


Figure 5: Training and validation loss across 5 epochs.

Table of Contents

- ① Introduction
- ② Related Work
- ③ Method
- ④ Fine-tuning the model
- ⑤ Playlist Generation
- ⑥ Results
- ⑦ Conclusion and Future Work

Playlist Generation: Overview

Goal: Recommend tracks matching the semantic theme of a playlist title.

- ① **Dataset embeddings:** Compute all playlists titles embeddings with the fine-tuned model.

⁶Tan Thongtan and Tanasanee Phienthrakul (July 2019). "Sentiment Classification Using Document Embeddings Trained with Cosine Similarity". In: *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics: Student Research Workshop*. Florence, Italy: Association for Computational Linguistics, pp. 407–414.

Playlist Generation: Overview

Goal: Recommend tracks matching the semantic theme of a playlist title.

- ① **Dataset embeddings:** Compute all playlists titles embeddings with the fine-tuned model.
- ② **User Input:** The user provides a playlist name.

⁶Tan Thongtan and Tanasanee Phienthrakul (July 2019). "Sentiment Classification Using Document Embeddings Trained with Cosine Similarity". In: *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics: Student Research Workshop*. Florence, Italy: Association for Computational Linguistics, pp. 407–414.

Playlist Generation: Overview

Goal: Recommend tracks matching the semantic theme of a playlist title.

- ① **Dataset embeddings:** Compute all playlists titles embeddings with the fine-tuned model.
- ② **User Input:** The user provides a playlist name.
- ③ **Input embedding:** The fine-tuned model converts the input title into a fixed-size embedding.

⁶Tan Thongtan and Tanasanee Phienthrakul (July 2019). "Sentiment Classification Using Document Embeddings Trained with Cosine Similarity". In: *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics: Student Research Workshop*. Florence, Italy: Association for Computational Linguistics, pp. 407–414.

Playlist Generation: Overview

Goal: Recommend tracks matching the semantic theme of a playlist title.

- ① **Dataset embeddings:** Compute all playlists titles embeddings with the fine-tuned model.
- ② **User Input:** The user provides a playlist name.
- ③ **Input embedding:** The fine-tuned model converts the input title into a fixed-size embedding.
- ④ **Similarity:** Compute cosine similarities⁶ between the input embedding and embeddings of all existing playlists.

⁶ Tan Thongtan and Tanasanee Phientrakul (July 2019). "Sentiment Classification Using Document Embeddings Trained with Cosine Similarity". In: *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics: Student Research Workshop*. Florence, Italy: Association for Computational Linguistics, pp. 407–414.

Playlist Generation: Overview

Goal: Recommend tracks matching the semantic theme of a playlist title.

- ① **Dataset embeddings:** Compute all playlists titles embeddings with the fine-tuned model.
- ② **User Input:** The user provides a playlist name.
- ③ **Input embedding:** The fine-tuned model converts the input title into a fixed-size embedding.
- ④ **Similarity:** Compute cosine similarities⁶ between the input embedding and embeddings of all existing playlists.

Why cosine similarity?

$$\text{cosine_similarity}(\mathbf{e}_a, \mathbf{e}_b) = \frac{\mathbf{e}_a \cdot \mathbf{e}_b}{\|\mathbf{e}_a\| \|\mathbf{e}_b\|}$$

- It measures the **angle** between two vectors, giving a sense of their directional similarity rather than magnitude.
- Well-suited for comparing high-dimensional embeddings.

⁶Tan Thongtan and Tanasanee Phienthrakul (July 2019). "Sentiment Classification Using Document Embeddings Trained with Cosine Similarity". In: *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics: Student Research Workshop*. Florence, Italy: Association for Computational Linguistics, pp. 407–414.

Playlist Generation: Voting system

- ① **Selecting Similar Playlists:** Identify k playlists with the highest cosine similarity scores.
- ② **Track Extraction:** Extract all tracks from these similar playlists.
- ③ **Voting Mechanism:** The N most frequently occurring songs are recommended.

How to select k and N

- Empirically, ensuring a balance between sufficient amount of tracks and high similarity scores.
- $k = 50$ (need at least 13).
- $N = 66$ (average number of tracks per playlist).

Table of Contents

① Introduction

② Related Work

③ Method

④ Fine-tuning the model

⑤ Playlist Generation

⑥ Results

⑦ Conclusion and Future Work

Results: Quantitative metrics

Quantitative Metrics⁷:

- **Precision@N:** Proportion of recommended tracks (S) that are relevant (R).

$$\text{precision}@N = \frac{\text{len}(R \cap S)}{\text{len}(S)} \quad (1)$$

- **Recall@N:** Proportion of relevant tracks that were successfully recommended.

$$\text{recall}@N = \frac{\text{len}(R \cap S)}{\text{len}(R)} \quad (2)$$

- **MRR@N (Mean Reciprocal Rank):** Focuses on the rank of the first relevant track.
 - A higher MRR indicates that the first correct recommendation appears earlier.

$$MRR@N = \frac{1}{\text{Rank of the first relevant song in the recommended songs}} \quad (3)$$

⁷ Guy Shani and Asela Gunawardana (2011). "Evaluating Recommendation Systems". In: *Recommender Systems Handbook*. Ed. by Francesco Ricci et al. Boston, MA: Springer US, pp. 257–297.

Results: Qualitative metrics

Qualitative Evaluation:

- Human assessment of thematic coherence and appeal.
- Evaluate each track to give the playlist an overall qualitative score.
- Important because numerical metrics may not fully capture user satisfaction.

Examples

For a playlist named "Rock classics"

- Highway to Hell, AC/DC
- Smells Like Teen Spirit, Nirvana
- It's my life, Bon Jovi
- My heart will go on, Céline Dion

Qualitative score 75%

Results: Experiment

Experiment Setup:

- Using test set playlists.
- Final playlist: 66 tracks.
- Three models were compared:
 - **Fine-tuned sentence-transformers** (expected to capture refined semantics)
 - **Pre-trained sentence-transformers** (baseline)
 - **Llama3.1** (LLM in zero-shot mode) (see 29).
- Evaluated on a test set of 22 carefully selected playlists.

Playlists selection criteria

- ≥ 10 tracks
- Clear thematic alignment between playlist titles and their tracks.
- A mix of both niche (e.g., 'hawaii') and more broadly themed playlists (e.g., 'summer').

LLM Zero-shot vs. LLM Multi-shot

LLM Zero-shot:

- Uses only an instruction prompt without examples, in a simple format.
- *"Give me the 66 most relevant songs that were released before October 2017 for a playlist named input. Do a list of songs using the format (song name, artist), without any other text formatting."*
- Simpler and faster, but may lack precision.

LLM Multi-shot (using ChatGPT 4o):

- **First idea:**
 - Provide the model with the list of songs from our most similar playlists.
 - **Problem:** would add bias.
- **Adjustment:** Provides the model with the dataset to yield context-aware outputs.

Comparison

Showing no relevant improvement.

Results: Quantitative model comparison

Model	Fine-tuned model	Pre-trained model	llama3.1
Precision@10	0.0586	0.0586	0.0008
Recall@10	0.1295	0.1295	0.0119
MRR@10	0.2189	0.2189	0.0109

Table 1: Comparison of quantitative scores between 3 models.

Key Points:

- The fine-tuning did not yield an improvement over the pre-trained model (same model).
- Hypothesis:** the weights were not correctly updated.
- Llama3.1's low scores suggest that our similarity-based voting approach is more relevant.

Results: Qualitative model comparison

Model	Fine-tuned model	Pre-trained model	llama3.1
Qualitative score	9.33	9.33	8.93

Table 2: Comparison of qualitative scores between 3 models.

Observations:

- Even when quantitative metrics were low, the recommended tracks often matched the theme.

Conclusion:

- Quantitative metrics alone do not capture the subjective quality of recommendations.
- Combining both evaluations provides a comprehensive view of system performance.

Table of Contents

① Introduction

② Related Work

③ Method

④ Fine-tuning the model

⑤ Playlist Generation

⑥ Results

⑦ Conclusion and Future Work

Conclusion

Summary:

- Developed a pipeline combining semantic clustering, transformer fine-tuning, and a similarity-based voting system for playlist continuation.
- Leveraged Sentence-BERT to capture the semantic nuances of playlist titles.
- Compared multiple models (fine-tuned Sentence-BERT, pre-trained Sentence-BERT, and LLM zero-shot)

What did not work:

- Fine-tuning a model.
- Observing results of our whole pipeline.

Future Work

Next Steps:

- **Resolve the fine-tuning issue** to ensure that model weights are effectively updated.
- **Explore alternative clustering techniques** to reduce “miscellaneous” clusters.
- **Optimize** the recommendation pipeline (e.g., experiment with different hyperparameters in the voting system).
- **Incorporate user feedback** for better qualitative scores.

What We Learned

Key Takeaways:

- **NLP fundamentals:** especially embeddings and BERT-based models.
- **Fine-tuning a model:** Work with LLMs to tailor them for specific tasks.
- **Evaluating recommender systems:** Precision@N and Recall@N.
- **Research-oriented approach:** combined practical coding with theoretical approach.
- SSH server and GitHub for code storage.

Demo

Live demo to show the final recommender system

Choose a playlist title.

Thank You!

Do you have any questions?

Bibliography I

-  Burke, Robin, Alexander Felfernig, and Mehmet H Göker (2011). "Recommender systems: An overview". In: *Ai Magazine* 32.3, pp. 13–18.
-  Chen, Ching-Wei et al. (2018). "Recsys challenge 2018: Automatic music playlist continuation". In: *Proceedings of the 12th ACM Conference on Recommender Systems*, pp. 527–528.
-  Espejel, Omar (2022). "Train and Fine-Tune Sentence Transformers Models". In: *Hugging Face*.
-  Li, Bohan et al. (2020). "On the sentence embeddings from pre-trained language models". In: *arXiv preprint arXiv:2011.05864*.
-  Monti, Diego et al. (2018). "An ensemble approach of recurrent neural networks using pre-trained embeddings for playlist completion". In: *Proceedings of the ACM Recommender Systems Challenge 2018*, pp. 1–6.

Bibliography II

-  Reimers, Nils and Iryna Gurevych (Nov. 2019). "Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks". In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing*. Association for Computational Linguistics, pp. 3982–3992.
-  Shani, Guy and Asela Gunawardana (2011). "Evaluating Recommendation Systems". In: *Recommender Systems Handbook*. Ed. by Francesco Ricci et al. Boston, MA: Springer US, pp. 257–297.
-  Sun, Chi et al. (2019). "How to fine-tune bert for text classification?" In: *Chinese computational linguistics: 18th China national conference, CCL 2019, Kunming, China, October 18–20, 2019, proceedings 18*. Springer, pp. 194–206.

Bibliography III

-  Thongtan, Tan and Tanasanee Phienthrakul (July 2019). "Sentiment Classification Using Document Embeddings Trained with Cosine Similarity". In: *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics: Student Research Workshop*. Florence, Italy: Association for Computational Linguistics, pp. 407–414.
-  Zhou, Yichu and Vivek Srikumar (2021). "A Closer Look at How Fine-tuning Changes BERT". In: *CoRR* abs/2106.14282. arXiv: 2106.14282. URL: <https://arxiv.org/abs/2106.14282>.